

ML project code-Copy1

March 10, 2022

```
[2]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt    # library providing tools for plotting data
from sklearn.preprocessing import PolynomialFeatures
import seaborn as sns    #data visualization library
from sklearn.linear_model import LinearRegression, HuberRegressor    # classes
    ↳ providing Linear Regression with ordinary squared error loss and Huber loss,
    ↳ respectively
from sklearn.metrics import mean_squared_error, mean_absolute_error
from sklearn.model_selection import train_test_split, KFold
```

```
[3]: data = pd.read_csv('Athlete_data.csv', sep = ';')
data.head(5)
```

```
[3]:
```

	year	400m hurdles	200 meters
0	2021	56.72	24.14
1	2021	56.94	24.88
2	2021	58.72	25.11
3	2021	58.92	25.36
4	2021	59.07	25.63

```
[4]: data2 = data.drop(['year'],axis = 1)

data2.columns = ['400mh','200m']

data2.head(5)
```

```
[4]:
```

	400mh	200m
0	56.72	24.14
1	56.94	24.88
2	58.72	25.11
3	58.92	25.36
4	59.07	25.63

```
[5]: x1 = data2['200m'].to_numpy()
X1 = x1.reshape(-1,1)
```

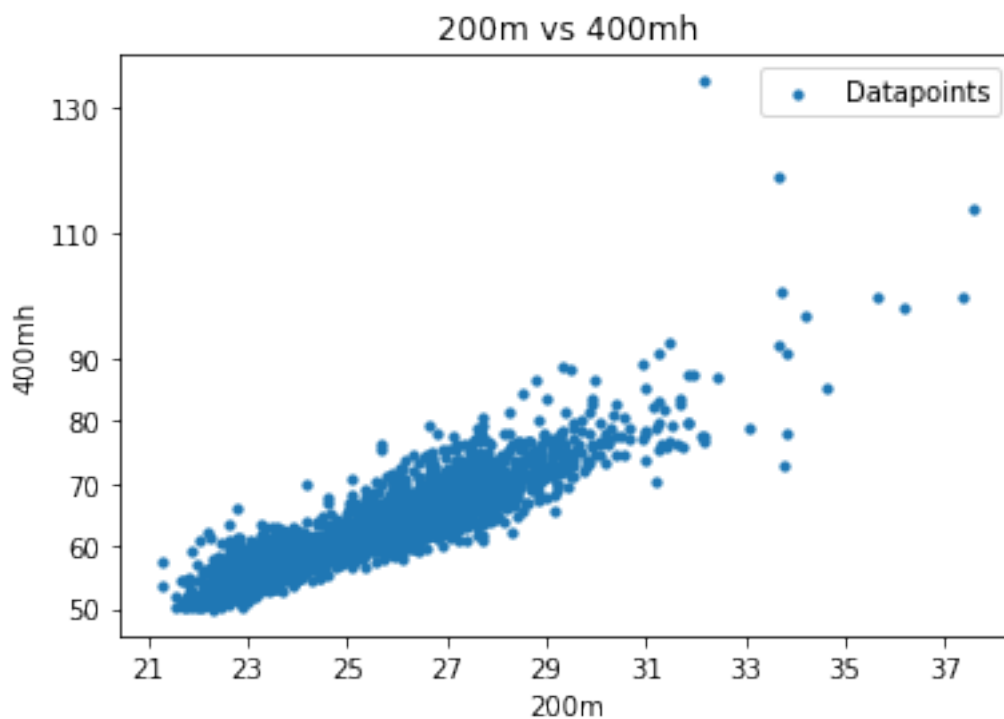
```
y1 = data2['400mh'].to_numpy()
```

```
y1.size
```

```
X1.shape
```

```
[5]: (2735, 1)
```

```
[6]: plt.scatter(X1,y1, s=10, label = "Datapoints");  
plt.xlabel("200m")  
plt.xticks([21,23,25,27,29,31,33,35,37])  
plt.ylabel("400mh")  
plt.yticks([50,60,70,80,90,110, 130])  
plt.title("200m vs 400mh")  
plt.legend(loc="best")  
  
plt.show()
```



```
[7]: trainingset_size = [0.4, 0.5, 0.6, 0.7]    # set the different sizes of training set  
  
for i in range(len(trainingset_size)):    # use for-loop to fit linear regression models with different sizes of training set
```

```

index = np.arange(int(len(X1)*trainingset_size[i]))
print("\nNumber of datapoints in this subset: ",len(index))

X_sub = X1[index]      # obtain a subset
y_sub = y1[index]

# Calculate training error of Huber model
hmodel = HuberRegressor()
hmodel.fit(X_sub,y_sub)
y_pred = hmodel.predict(X_sub)
tr_error = mean_squared_error(y_sub, y_pred)

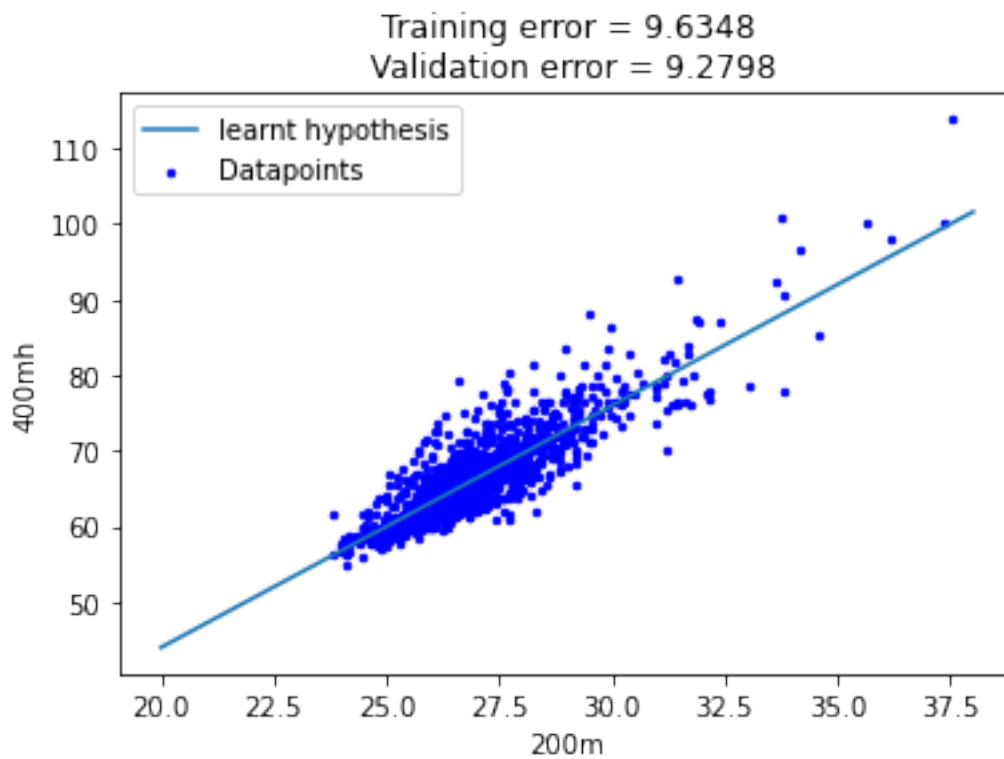
#Calculate validation error
val_index = np.arange(int(len(X1)*(1-trainingset_size[i])))
X_val_sub = X1[val_index]
y_val_sub = y1[val_index]

y_pred_val_sub = hmodel.predict(X_val_sub)
val_error_sub = mean_squared_error(y_val_sub,y_pred_val_sub)


X_fit = np.linspace(20, 38, 100)      # generate samples
plt.plot(X_fit, hmodel.predict((X_fit.reshape(-1, 1))), label="learnt_
→hypothesis")      # plot the linear regression model
plt.scatter(X_sub, y_sub, color = 'b', s=6, label = 'Datapoints')      # plot_
→a scatter plot of y(200m) vs. X(400mh)
plt.xlabel('200m')      # set the label for the x/y-axis
plt.ylabel('400mh')
plt.legend(loc="best")      # set the location of the legend
plt.title(f'Training error = {tr_error:.5}\n Validation error =_
→{val_error_sub:.5}')      # set the title
plt.show()      # show the plot

```


Number of datapoints in this subset: 1094




Number of datapoints in this subset: 1367



Number of datapoints in this subset: 1641



Number of datapoints in this subset: 1914



```
[8]: ## Splitting data with 8:2 ratio
X_train, X_test, y_train, y_test = train_test_split(X1,y1, test_size = 0.2,
↳random_state = 42)
```

```
[ ]:
```

```
[9]: ## Training set evaluated/trained using k-fold cross validation

cv = KFold(n_splits = 10, shuffle = True, random_state = 42)
val_errors2 = []
tr_errors2 = []
# Iterate through the indices of train and validation (iteration through each
↳split of 20)
for train_index, val_index in cv.split(y_train):

    X_train2, X_val2 = X_train[train_index], X_train[val_index]
    y_train2, y_val2 = y_train[train_index], y_train[val_index]

    hmodel3 = HuberRegressor()
```

```

hmodel3.fit(X_train2, y_train2)# apply Huber regression to these new
↪features and labels
y_pred_train2 = hmodel3.predict(X_train2)
tr_error3 = mean_squared_error(y_train2, y_pred_train2)
print('Training error = {:.5}'.format(tr_error3))
tr_errors2.append(tr_error3)

#predict values for the validation data using the linear model
#calculate the validation error

y_pred_val2 = hmodel3.predict(X_val2)
val_error2 = mean_squared_error(y_val2, y_pred_val2)

print('Validation error= {:.5}'.format(val_error2))
val_errors2.append(val_error2)

```

```

Training error = 10.617
Validation error= 8.2603
Training error = 10.474
Validation error= 9.8118
Training error = 9.01
Validation error= 23.182
Training error = 10.742
Validation error= 7.8405
Training error = 10.652
Validation error= 8.1054
Training error = 10.434
Validation error= 10.627
Training error = 10.073
Validation error= 13.853
Training error = 10.869
Validation error= 6.23
Training error = 10.554
Validation error= 9.1736
Training error = 10.775
Validation error= 7.276

```

```

[10]: #Calculate average validation error
sum = 0
for i in val_errors2:
    sum = sum + i

avg_val_error = sum/(len(val_errors2))
print('Average validation error= {:.5}'.format(avg_val_error))

```

```

Average validation error= 10.436

```



```
[11]: #Calculate average training error
      sum2 = 0
      for i in tr_errors2:
          sum2 = sum2 +i

      avg_train_error = sum2/(len(tr_errors2))
      print('Average training error = {:.5}'.format(avg_train_error))
```

Average training error = 10.42

```
[ ]:
```